

```
//*****  
/* Copyright (C) 2018 International Business Machines Corporation  
/* All Rights Reserved  
//*****
```

How to connect to the HBASE database of Analytics Engine on IBM Cloud via JDBC toolkit

This document describes a step by step instruction to connect to a HBASE database on IBM cloud via **JDBC** toolkit.

1 - Download the phoenix jdbc drivers from:

<https://github.com/IBM-Cloud/IBM-Analytics-Engine/tree/master/jdbcsamples/TestPhoenix/lib>

and copy them in the **opt** directory in your project workspace:

```
JDBCHbase/opt/phoenix-4.9.0-HBase-1.1-client.jar  
JDBCHbase/opt/phoenix-queryserver-4.9.0-HBase-1.1.jar  
JDBCHbase/opt/phoenix-tracing-webapp-4.9.0-HBase-1.1.jar  
JDBCHbase/opt/phoenix-4.9.0-HBase-1.1-queryserver.jar  
JDBCHbase/opt/phoenix-queryserver-client-4.9.0-HBase-1.1.jar
```

2 - Create an Analytics Engine service on IBM Cloud and create a service credential

Create a Analytics Engine service on IBM cloud.

<https://console.bluemix.net/catalog/?search=Analytics%20Engine>

IBM Analytics Engine documentation

<https://console.bluemix.net/docs/services/AnalyticsEngine/index.html#introduction>

Create a service credential for Analytics Engine service on IBM cloud.

The following sample shows an Analytics Engine service credentials:

```
{  
  "apikey": "xyxyxyxyxyxyxyxyxyxyxyxyxyxyxy",  
  "cluster": {  
    "cluster_id": "20180404-125209-123-VNhbnQRZ",  
    "password": "IAEPASSWORD",  
    "service_endpoints": {  
      "ambari_console": "https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:9443",  
      "hive_jdbc": "jdbc:hive2://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/?ssl=t",  
      "livy": "https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/gateway/default/",  
      "notebook_gateway": "https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/gate",  
      "notebook_gateway_websocket": "wss://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8",  
      "oozie_rest": "https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/gateway/de",  
      "phoenix_jdbc": "jdbc:phoenix:thin:url=https://chs-nxr-123-mn001.bi.services.us-south.blu
```

```

    "spark_history_server": "https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/",
    "spark_sql": "jdbc:hive2://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/?ssl=t",
    "ssh": "ssh clsadmin@chs-nxr-123-mn003.bi.services.us-south.bluemix.net",
    "webhdfs": "https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:8443/gateway/default",
  },

  "user": "clsadmin"
}
},

```

For a JDBC connection to HBASE we need from Analytics Engine service credentials 3 parameters:

The jdbcurl in this sample DB2 service credentials is:

```
"jdbcurl": "jdbc:phoenix:thin:url=https://chs-nxr-123-mn001.bi.services.us-south.bluemix.net:
```

Copy "jdbcurl" "user" and "password" from IBM Analytics Engine service credential and put it in the following spl file

2 - Create a Streams application

Create a new project **JDBCHbase** in your workspace:

```
~/workspace/JDBCHbase/application/JDBCHbase.spl
~/workspace/JDBCHbase/Makefile
```

JDBCHbase.spl

```

// *****
// * Copyright (C) 2018 International Business Machines Corporation
// * All Rights Reserved
// *****
// JDBCRUN_HBASE.spl demonstrates how to connect to a HBASE database
// on IBM Analytics Engine via streams toolkit com.ibm.streamsx.jdbc
// and Apache Phoenix JDBC
// It demonstrates also how to:
// create a table
// insert data into table
// select data from table
// and drop a table
// via JDBCRun operator.
//
// To connect to database, the following parameters need to be specified:
// jdbcDriverLib : the jdbc driver library (download the jdbc driver and store it in opt folder
// e.g. opt/db2jcc4.jar)
// jdbcClassName : the class name for jdbc driver (e.g. com.ibm.db2.jcc.DB2Driver)
// jdbcUrl       : This parameter specifies the database url that JDBC driver uses
// to connect to a database and it must have exactly one value of type rstring.
// The syntax of jdbc url is specified by database vendors.
// For example, jdbc:db2://<server>:<port>:<database>the database URL.
// and For Phoenix JDBC "jdbc:phoenix:thin:url=https://<database-host-name>:<server:port>/<database>
// jdbcUser      : the database user on whose behalf the connection is being made.
// jdbcPassword  : the database user's password.
//
// This sample runs with streams JDBC toolkit version 1.3 or higher
// https://github.com/IBMStreams/streamsx.jdbc
// and with phoenix jdbc driver.
// https://github.com/IBM-Cloud/IBM-Analytics-Engine/tree/master/jdbcsamples/TestPhoenix/lib
// or direct from https://phoenix.apache.org/index.html

```

```
// Download and copy the following jar libraries and copy them in opt directory
// phoenix-4.9.0-HBase-1.1-client.jar
// phoenix-queryserver-4.9.0-HBase-1.1.jar
// phoenix-tracing-webapp-4.9.0-HBase-1.1.jar
// phoenix-4.9.0-HBase-1.1-queryserver.jar
// phoenix-queryserver-client-4.9.0-HBase-1.1.jar
//
// *****
```

```
namespace application;
use com.ibm.streamsx.jdbc::* ;
use com.ibm.streamsx.jdbc.types::* ;
```

```
composite JDBCBase
```

```
{
    param
    ⚡expression<rstring> $jdbcDriverLib⚡ getSubmissionTimeValue("jdbcDriverLib", "opt/phoenix-core
    ⚡expression<rstring> $jdbcClassName⚡ getSubmissionTimeValue("jdbcClassName", "org.apache.phoen
    ⚡expression<rstring> $jdbcUrl ⚡⚡getSubmissionTimeValue("jdbcUrl", "jdbc:phoenix:thin:url=https
    ⚡expression<rstring> $jdbcUser ⚡⚡getSubmissionTimeValue("jdbcUser", "clsadmin");
    ⚡expression<rstring> $jdbcPassword ⚡ getSubmissionTimeValue("jdbcPassword", "IAEPASSWORD");
```

```
    ⚡expression<rstring> $createSql ⚡⚡"CREATE TABLE IF NOT EXISTS JDBCRUN_HBASE (id bigint ,m.fname
    ⚡expression<rstring> $selectSql ⚡⚡"SELECT * FROM JDBCRUN_HBASE";
    ⚡expression<rstring> $dropSql ⚡⚡"DROP TABLE IF EXISTS JDBCRUN_HBASE";
    ⚡expression<rstring> $insertSql ⚡⚡"UPSERT INTO JDBCRUN_HBASE VALUES( ";
```

```
    type
    ⚡InsertSchema = int32 ID, rstring FNAME, rstring LNAME;
    ⚡ResultsSchema = int32 ID, rstring FNAME, rstring LNAME;
```

```
graph
```

```
    /**
    * The pulse is a Beacon operator that generates counter.
    */
```

```
⚡Stream<int32 counter> pulse = Beacon() {
    logic
    ⚡State: mutable int32 i = 0;
```

```
param
    ⚡initDelay : 2.0;
    ⚡iterations : 11u ;
    ⚡period⚡⚡3.0;
    output
    ⚡pulse : counter = i++;
}
```

```
    /**
    * genStatement is Custom operator that generates sql statements.
    */
```

```
⚡Stream<rstring statement> genStatement = Custom(pulse)
{
    logic
```

```
⚡State: mutable rstring ⚡stmt = "UPSERT INTO JDBCRUN_HBASE values( 0, 'FNAME0', 'LNAME0')";⚡⚡⚡⚡
```

```
onTuple pulse :
{
    ⚡⚡
    ⚡if⚡ ( counter == 0)
    {
    ⚡stmt = $dropSql;⚡⚡⚡⚡
    }
    else if ( counter == 1)
    {

```

```

stmt = $createSql;
}
else if ( counter == 10)
{
stmt = $selectSql;
}
else
{
stmt = $insertSql + (rstring)counter + ", 'FNAME" + (rstring)counter + "' , 'LNAME" + (rstri:
}
}
}

println( (rstring)counter + " SQL Statement = " + stmt );
submit({ statement = stmt }, genStatement) ;
}
}

/**
 * runSql is a JDBCRun operator.
 * It runs a user-defined SQL statement that is based on an input tuple.
 */
Stream<insertSchema> runSql = JDBCRun(genStatement)
{
param
jdbcDriverLib: $jdbcDriverLib ;
jdbcClassName: $jdbcClassName ;
jdbcUrl: $jdbcUrl ;
jdbcUser: $jdbcUser ;
jdbcPassword: $jdbcPassword;
statementAttr: statement ;
sqlFailureAction : "log";
}

/**
 * printResults is a Custom operator that prints the sql results.
 */
() as printResults = Custom(runSql)
{
logic
onTuple runSql :
{
if(ID > 0)
{
println((rstring) ID + "," + FNAME + "," + LNAME);
}
}
}
}

```

4 - Make the SPL application

To create this SPL application the new version of JDBC toolkit (1.3.0) is required.

Download and copy the latest version of streamsx.jdbc in your workspace.

<https://github.com/IBMStreams/streamsx.jdbc>

The Makefile makes also the toolkit.

```

#####
# Copyright (C)2014, 2018 International Business Machines Corporation and

```

```
# others. All Rights Reserved.
#####

.PHONY: all clean

//SPLC_FLAGS = -t $(STREAMS_INSTALL)/toolkits/com.ibm.streamsx.jdbc --data-directory data
SPLC_FLAGS = -t ../streamsx.jdbc/com.ibm.streamsx.jdbc --data-directory data

SPLC = $(STREAMS_INSTALL)/bin/sc
SPL_CMD_ARGS ?=
SPL_COMPINAME=JDBCHbase
SPL_MAIN_COMPOSITE1 = application::$(SPL_COMPINAME)
BUILD_OUTPUT_DIR = output

all: clean
    $(SPLC) $(SPLC_FLAGS) -M $(SPL_MAIN_COMPOSITE1) --output-dir ./${BUILD_OUTPUT_DIR}

clean:
    $(SPLC) $(SPLC_FLAGS) -C -M $(SPL_MAIN_COMPOSITE1) --output-dir output
    -rm -rf toolkit.xml
```

Be aware of tabs in Makefile

5 - Run the SPL application

Change the database credentials in SPL file with your IBM HBASE database credentials and run

```
$> make
```

Start the application with

```
$> output/bin/standalone
```

Or you can submit the job on your local Streams server with:

```
$ streamtool submitjob output/application.JDBCHbase.sab
```

6 - Submit the spl application on IBM Streams Cloud

Create a Streaming Analytics on IBM Cloud

<https://console.bluemix.net/catalog/?search=streams>

Start the service

Lunch the application

It starts the IBM Streams console.

Now it is possible here to submit a SAB file as job

The SAB file is located in your project output directory:

7 - Check the result on hbase server

To check the result login on IAE server and check the contain of table with hbase shell.

```
$> ssh clsadmin@<your-IAE-ip-address>
$ > hbase shell
```

```
hbase(main):001:0> list
JDBCRUN_HBASE
SYSTEM.CATALOG
SYSTEM.FUNCTION
SYSTEM.SEQUENCE
SYSTEM.STATS    □
...

hbase(main):001:0> scan 'JDBCRUN_HBASE'
ROW                                COLUMN+CELL
\x80\x00\x00\x00\x00\x00\x00\x02  column=M:FNAME, timestamp=1523608569941, value=FNAME2
\x80\x00\x00\x00\x00\x00\x00\x02  column=M:LNAME, timestamp=1523608569941, value=LNAME2
\x80\x00\x00\x00\x00\x00\x00\x02  column=M:_0, timestamp=1523608569941, value=x
.....
```

For more details about the **hbase shell** check this link:

<https://github.com/IBMStreams/streamsx.jdbc>